

Assignment No. 6

Problem Statement: Design and implement the **Apriori Algorithm** to find frequent itemsets and generate association rules from a given transactional dataset.

Objective:

1. To understand the concept of frequent itemset mining and association rule learning.
2. To implement the Apriori algorithm using an appropriate data mining library.
3. To analyze the output of the algorithm and interpret association rules (support, confidence, lift).
4. To identify the most relevant patterns within a dataset based on minimum support and confidence thresholds.

Prerequisite :

1. **Basic Programming Concepts** – Python, loops, data structures.
2. **Set Theory & Probability** – Understanding of sets, unions, intersections, probability
3. **Data Mining Concepts** – Understanding of support, confidence, lift, itemsets.
4. **Libraries** – Basic understanding of pandas, mlxtend, or scikit-learn.

Theory :

The Apriori algorithm is a fundamental and widely used technique in the field of data mining, particularly for the task of **association rule learning**. It is designed to extract meaningful patterns from large transactional datasets by identifying frequent combinations of items (known as **frequent itemsets**) and using these combinations to infer **association rules**. This algorithm is most commonly associated with **market basket analysis**, where the objective is to understand customer purchasing behavior by analyzing the items bought together.

Apriori is based on the **anti-monotonicity property** of support, also called the **Apriori Principle**, which states:

"If an itemset is frequent, then all of its subsets must also be frequent."

This property allows Apriori to efficiently prune the search space and avoid evaluating itemsets that cannot possibly be frequent, significantly reducing computational overhead in large databases.

Key Terminologies

1. Itemset

An itemset refers to a collection of one or more items. These items are typically part of transactions. For example, in a supermarket, a transaction might consist of items like milk, bread, and butter. A 1-itemset contains one item (e.g., {milk}), a 2-itemset contains two items (e.g., {milk, bread}), and so on.

2. Support

Support is a measure of how frequently an itemset appears in the dataset. It reflects the proportion of transactions that include a particular itemset. Mathematically, it is defined as:

$$\text{Support}(X) = \frac{|\{T \in D : X \subseteq T\}|}{|D|}$$

Where:

1. XXX is an itemset,
2. TTT is a transaction,
3. DDD is the database of all transactions.

A high support indicates that the itemset is common and may be useful in discovering stable patterns.

3. Confidence

Confidence is a measure of the reliability of an association rule. It indicates the probability that a transaction containing itemset X also contains itemset Y. In rule notation $X \rightarrow Y$ confidence is given by:

$$\text{Confidence}(X \rightarrow Y) = \frac{\text{Support}(X \cup Y)}{\text{Support}(X)}$$

4. Lift

Lift measures the strength of a rule compared to the random co-occurrence of the antecedent and consequent. It quantifies how much more likely items in Y are purchased when X is purchased, compared to when they are purchased independently. It is defined as:

$$\text{Lift}(X \rightarrow Y) = \frac{\text{Confidence}(X \rightarrow Y)}{\text{Support}(Y)}$$

- **If lift = 1: X and Y are independent.**
- **If lift > 1: X and Y are positively correlated.**
- **If lift < 1: X and Y are negatively correlated.**

5. **Frequent Itemset**

A frequent itemset is any group of items that appears together in at least a specified minimum number of transactions (as determined by the support threshold).

6. **Association Rule**

An association rule is an implication of the form $X \rightarrow Y$, where X and Y are disjoint itemsets. The rule suggests that if X occurs in a transaction, then Y is likely to also occur.

Working of the Apriori Algorithm

1. Finding Frequent Itemsets

This step aims to discover all item combinations that occur in the dataset with a frequency greater than or equal to the user-defined minimum support threshold.

1. **Initialization ($k = 1$):**
Identify all frequent 1-itemsets by counting their individual occurrences in the dataset.
2. **Iteration ($k \geq 2$):**
For each level k , generate candidate k -itemsets by joining frequent $(k-1)$ -itemsets with each other.
3. **Pruning:**
Remove those candidate itemsets where any $(k-1)$ -subset is not frequent. This is based on the Apriori property.
4. **Support Counting:**
Count the support of each candidate k -itemset and remove those below the threshold.
5. **Repeat**
Continue the process until no new frequent itemsets are found.

2. **Generating Association Rules**

Once frequent itemsets are identified, the next step is to generate rules from them.

1. For each frequent itemset L, generate all non-empty subsets.
2. For each subset S of L, construct the rule $S \rightarrow (L-S)$
3. Compute confidence and lift for each rule.
4. Retain only those rules that meet the minimum confidence and lift thresholds.

For example, from frequent itemset {A,B,C} we can generate:

1. $A \rightarrow B, C$
2. $B \rightarrow A, C$

3. $C \rightarrow A, B$
4. and so on.

Optimizations

Several optimizations have been proposed to improve the efficiency of Apriori:

1. **Hash-Based Candidate Pruning:** Uses hash tables to reduce the number of candidate itemsets generated in the early stages.
2. **Transaction Reduction:** Removes transactions that do not contain any frequent itemsets, reducing the size of the dataset.
3. **Partitioning:** Divides the database into partitions and mines frequent itemsets locally before combining them globally.
4. **Sampling:** Selects a representative sample of the dataset to approximate frequent itemsets.

Applications

Apriori and its variants are used across a wide range of applications:

1. **Retail and Market Basket Analysis** – Identifying products that are frequently bought together.
2. **Healthcare** – Detecting co-occurrence of symptoms and treatments.
3. **E-commerce Recommendation Engines** – Suggesting products based on association rules.
4. **Fraud Detection** – Finding unusual combinations of transaction patterns.
5. **Web Usage Mining** – Analyzing user navigation patterns on websites.

Advantages

1. Conceptually simple and easy to implement.
2. Effectively uses the Apriori property to reduce the number of candidate itemsets.
3. Generates interpretable rules useful for strategic decisions.

Limitations

1. Generates a vast number of candidate itemsets when the dataset has a large number of items, leading to high memory and computational costs.
2. Requires multiple passes over the entire dataset, which is time-consuming for large databases.
3. Poor performance with low support thresholds.
4. Does not handle numeric or continuous data natively; requires discretization.

Code & Output

```
[1]: import numpy as np
import pandas as pd
```

```
[15]: df = pd.read_csv("C:/Users/dnyan/ML Assignments/Dataset/store_data.csv")
```

```
[16]: print(df.head())
```

```
      shrimp  almonds  avocado  vegetables mix green grapes \
0    burgers  meatballs      eggs      NaN      NaN
1    chutney      NaN      NaN      NaN      NaN
2    turkey  avocado      NaN      NaN      NaN
3  mineral water      milk  energy bar  whole wheat rice  green tea
4 low fat yogurt      NaN      NaN      NaN      NaN

  whole weat flour yams cottage cheese energy drink tomato juice \
0      NaN  NaN      NaN      NaN      NaN      NaN      NaN
1      NaN  NaN      NaN      NaN      NaN      NaN      NaN
2      NaN  NaN      NaN      NaN      NaN      NaN      NaN
3      NaN  NaN      NaN      NaN      NaN      NaN      NaN
4      NaN  NaN      NaN      NaN      NaN      NaN      NaN

  low fat yogurt green tea honey salad mineral water salmon antioxydant juice \
0      NaN      NaN  NaN  NaN      NaN      NaN      NaN      NaN
1      NaN      NaN  NaN  NaN      NaN      NaN      NaN      NaN
2      NaN      NaN  NaN  NaN      NaN      NaN      NaN      NaN
3      NaN      NaN  NaN  NaN      NaN      NaN      NaN      NaN
4      NaN      NaN  NaN  NaN      NaN      NaN      NaN      NaN
```

```
[17]: # shape of the data
df.shape
```

```
[17]: (7500, 20)
```

```
[18]: #data information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7500 entries, 0 to 7499
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   shrimp                7500 non-null  object
1   almonds               5746 non-null  object
2   avocado               4388 non-null  object
3   vegetables mix        3344 non-null  object
4   green grapes          2528 non-null  object
5   whole weat flour      1863 non-null  object
6   yams                  1368 non-null  object
7   cottage cheese        980 non-null   object
8   energy drink          653 non-null   object
9   tomato juice          394 non-null   object
10  low fat yogurt        255 non-null   object
11  green tea             153 non-null   object
12  honey                 86 non-null    object
13  salad                 46 non-null    object
14  mineral water         24 non-null    object
15  salmon                7 non-null     object
16  antioxydant juice     3 non-null     object
17  frozen smoothie       3 non-null     object
18  spinach               2 non-null     object
19  olive oil             0 non-null     float64
dtypes: float64(1), object(19)
memory usage: 1.1+ MB
```

```
df = df.drop(columns=['olive oil'])

# Convert each column into binary: 1 if item exists (non-null), 0 otherwise.
# We use .notnull() which returns True/False, then convert Boolean to integer (1/0)
df_binary = df.notnull().astype(int)

# Quick check of the preprocessed binary data
print("\nBinary Data Sample:")
print(df_binary.head())
```

Binary Data Sample:

	shrimp	almonds	avocado	vegetables mix	green grapes	whole wheat flour	\
0	1	1	1	0	0	0	
1	1	0	0	0	0	0	
2	1	1	0	0	0	0	
3	1	1	1	1	1	0	
4	1	0	0	0	0	0	

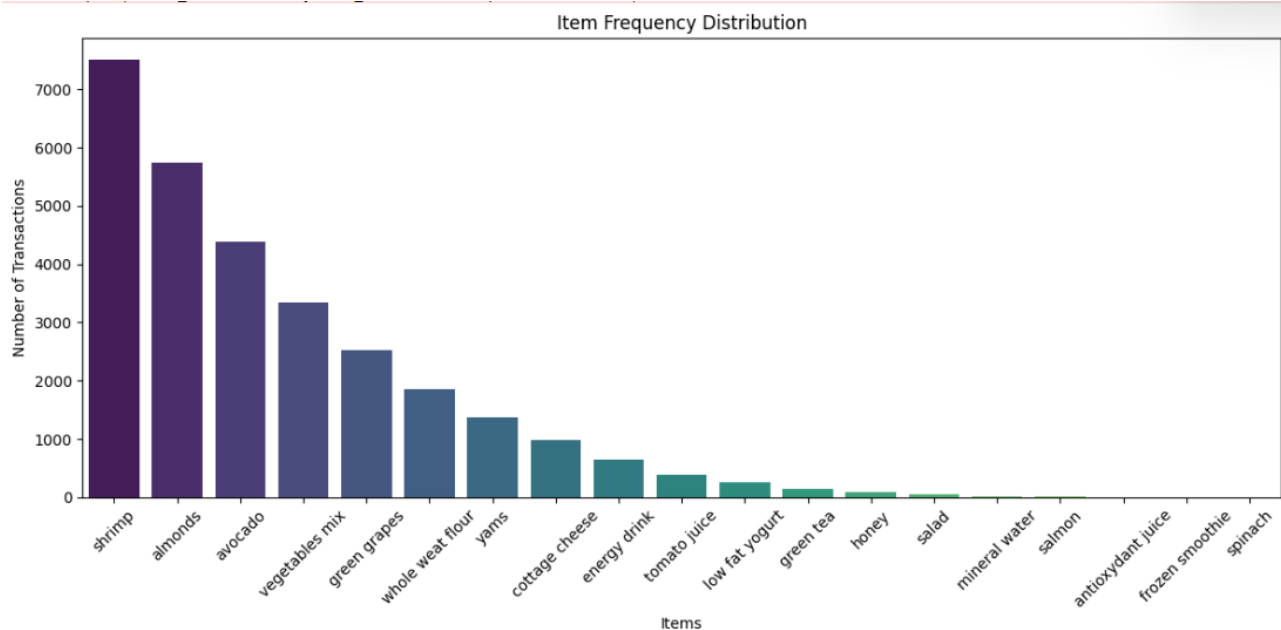
	yams	cottage cheese	energy drink	tomato juice	low fat yogurt	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	green tea	honey	salad	mineral water	salmon	antioxydant juice	\
0	0	0	0	0	0	0	
1	0	0	0	0	0	0	
2	0	0	0	0	0	0	
3	0	0	0	0	0	0	
4	0	0	0	0	0	0	

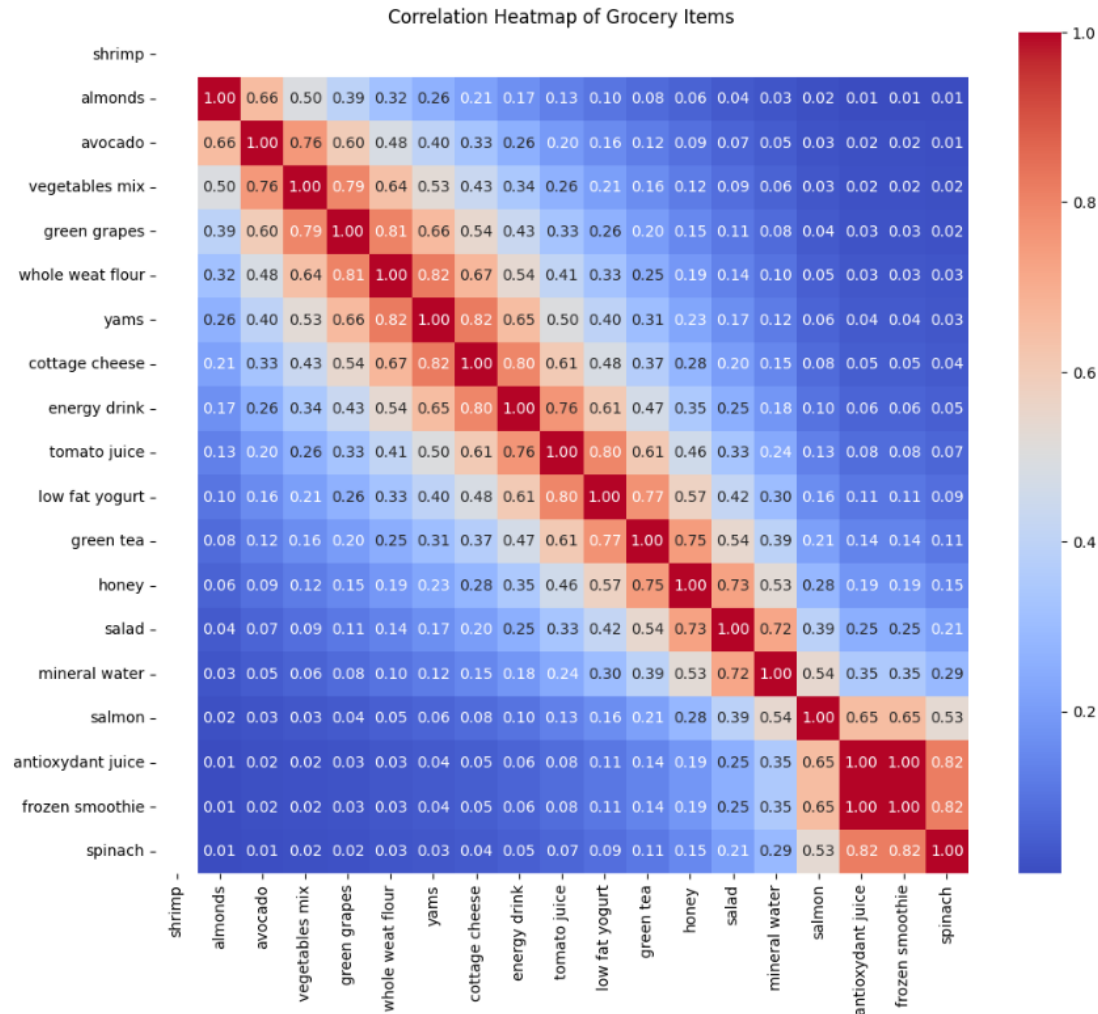
```
import matplotlib.pyplot as plt
import seaborn as sns

# Sum each column to count the number of transactions that include the item
item_counts = df_binary.sum().sort_values(ascending=False)

plt.figure(figsize=(12, 6))
sns.barplot(x=item_counts.index, y=item_counts.values, palette='viridis')
plt.title("Item Frequency Distribution")
plt.xlabel("Items")
plt.ylabel("Number of Transactions")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
plt.figure(figsize=(12, 10))
sns.heatmap(df_binary.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Heatmap of Grocery Items")
plt.show()
```



```
from mlxtend.frequent_patterns import apriori

# Calculate frequent itemsets with a minimum support threshold.
frequent_itemsets = apriori(df_binary, min_support=0.01, use_colnames=True)

print("Frequent Itemsets:")
print(frequent_itemsets.sort_values(by="support", ascending=False).head())
```

C:\Users\dnyan\AppData\Local\Programs\Python\Python312\Lib\site-packages\mlxtend\non-bool types result in worse computational performance and their support might


warnings.warn(

Frequent Itemsets:

	support	itemsets
0	1.000000	(shrimp)
13	0.766133	(shrimp, almonds)
1	0.766133	(almonds)
14	0.585067	(shrimp, avocado)
25	0.585067	(almonds, avocado)

```
plt.figure(figsize=(10, 6))
sns.scatterplot(x="support", y="confidence", size="lift",
               data=rules, sizes=(20, 200), alpha=0.7)
plt.title("Association Rules: Support vs Confidence")
plt.xlabel("Support")
plt.ylabel("Confidence")
plt.show()
```

C:\Users\dnyan\AppData\Local\Programs\Python\Python312\Lib\site-packages\IPython\core\pylabtools.py:170: UserWarning: Creating a figure with a size of (10, 6) may be slow with large amounts of data.
fig.canvas.print_figure(bytes_io, **kw)



Github : <https://github.com/dnyaneshwardhere/ML>

Conclusion:

The Apriori algorithm is a basic but powerful method for finding frequent itemsets and generating association rules in transactional data. It uses the principle that all subsets of a frequent itemset must also be frequent, which helps reduce unnecessary computations. While simple and easy to understand, it can be slow for large datasets. Despite this, it remains a valuable tool for learning and discovering useful patterns in fields like retail, healthcare, and web analytics.